The Factor Skyline: A Geometric Architecture for Multiplicative Number Theory

Allen Proxmire
March 2026

Abstract
The Factor Skyline introduces a two-dimensional geometric architecture in which each integer is represented as a column whose width is its least prime factor and whose height is the quotient after removing that factor. This simple construction reveals a deterministic sieve operating in physical space: each prime creates an independent coverage layer, activation occurs at prime squares, and primality emerges as the escape from all layers. From this mechanism arise the classical results of multiplicative number theory, including the Prime Number Theorem, the Dickman function, squarefree density, Möbius cancellation, divisor statistics, and the conservation law for escape height.

The skyline also unifies the statistical behavior of the integers. Shared and independent coverage layers generate the full hierarchy of correlations—escape, parity, branching, and spectral—while the interaction between the deterministic template and the thin escape layer produces the three-level randomness hierarchy. Information-theoretic analysis shows that the template removes most of the uncertainty in primality, leaving an irreducible core of approximately 0.26 bits per integer. This core is the parity barrier, the structural limit underlying the difficulty of the major open conjectures.

Beyond the integers, the skyline defines a universal architecture for multiplicative systems satisfying four minimal axioms. It identifies three universality classes, twelve universal structures, and a renormalization flow governed by the escape density. At the meta-level, the skyline exhibits a self-referential structure: the emergence hierarchy, dualities, computational complexity, and logical limits mirror the architecture of the integers themselves.

The Factor Skyline is therefore not merely a visualization of the integers. It is a unified geometric, statistical, informational, and meta-mathematical framework that reorganizes multiplicative number theory into a single coherent architecture.

1. The Big Idea
Prime numbers have always looked mysterious. They appear irregularly, resist prediction, and seem to flicker in and out of the number line with no obvious pattern. For centuries, mathematicians have described primes as "random," "chaotic," or "unpredictable," even though they arise from the most deterministic structure imaginable: the integers themselves.

The Factor Skyline begins from a simple question: What if primes only look random because we've been looking at the integers the wrong way? Instead of treating the number line as a flat, one-dimensional sequence, the skyline lifts each integer into a two-dimensional geometric object. Every number becomes a column whose width is its least prime factor and whose height is the quotient after dividing by that factor. When these columns are placed side by side, a new geometry emerges — one in which the mechanisms behind prime behavior become visible.

The central discovery is that primes are not random at all. They are the deterministic escape points of a geometric coverage process. Each prime p creates a "coverage layer" that eliminates a predictable fraction of the integers. These layers act independently, and the only numbers that survive all of them are the primes. In this view,

primality is not a mysterious property. It is simply the condition of having escaped every smaller prime's coverage.

This single mechanism — independent coverage layers and deterministic escape — turns out to explain nearly everything in multiplicative number theory. The thinning of primes, the distribution of smooth numbers, the behavior of the Möbius function, the structure of prime constellations, the sub-Poisson variance of prime counts, the Hardy–Littlewood constants, the GUE statistics of zeta zeros — all of these arise naturally from the skyline's geometry.

The skyline also reveals why primes appear random. Although the coverage process is fully deterministic, the independence of the layers creates statistical noise that mimics randomness. The skyline quantifies this illusion: most of the information about primality is already determined by the template created by small primes, and only a thin slice — about 0.26 bits per integer — remains genuinely uncertain. This irreducible uncertainty is the information-theoretic form of the parity barrier, the structural limit that underlies the difficulty of the major open conjectures.

The Factor Skyline is not a new number system and does not alter the integers. It is a new ontology — a new way of seeing — that reorganizes the familiar facts of number theory into a coherent geometric architecture. In this architecture, the classical results are not isolated theorems but manifestations of a single underlying structure. The skyline makes that structure explicit.

This section introduces the core idea: primes are deterministic escapes from independent coverage layers, and the entire landscape of multiplicative number theory can be understood as the geometry of this escape process.

2. The Architecture of the Skyline

The Factor Skyline begins with a simple but powerful shift: instead of treating each integer as a point on a line, we treat it as a geometric object. Every number becomes a column whose shape is determined entirely by its least prime factor. This single choice — using the least prime factor as the generator — produces a surprisingly rich geometry, one that reveals the hidden structure of the integers.

The skyline is built from five architectural ideas: width, height, activation, coverage, and escape. Together, they form a deterministic sieve that operates in physical space.

Width is the least prime factor of a number. A prime has width 1, the narrowest possible column. A composite number has width equal to its smallest prime divisor: 4 has width 2, 9 has width 3, 25 has width 5, and so on. This width is not just a label — it determines how much horizontal space the number occupies in the skyline.

Height is what remains after dividing the number by its least prime factor. A prime has height equal to itself. A composite number has height equal to the quotient after peeling off its smallest prime. This height captures how "deep" the number is in the factorization process.

When these columns are placed side by side, the skyline emerges. The horizontal coordinate of each number is the cumulative sum of the widths of all previous numbers. Primes contribute only one unit of width, so they appear as thin spires. Composites contribute wider blocks, creating broad plateaus. The skyline becomes a landscape of narrow escapes and wide coverage.

This geometry is not arbitrary. It encodes the action of the sieve directly into space. Each prime p introduces a coverage layer that eliminates a predictable fraction of the integers: exactly 1/p of all positions fall under the shadow of p. These layers act independently, and their combined effect is multiplicative. The fraction of integers that escape all primes up to p is the product of the escape fractions of each layer. This product is the escape density, and it governs the distribution of primes.

A remarkable feature of the skyline is the role of prime squares. A prime p does not become the least prime factor of any number until p^2. Below p^2, every multiple of p has a smaller prime factor. This means that p^2 is the moment when the width-p layer "activates" and begins to cover new territory. These activation points divide the skyline into epochs — intervals between consecutive prime squares where the coverage pattern is frozen. Within an epoch, the geometry is stable; only at the next prime square does a new layer enter the system.

The skyline also contains a periodic backbone called the template. The template is the repeating pattern of open and covered positions modulo the primorial p#, the product of all primes up to p. This pattern determines which positions are guaranteed composite and which remain open to escape. The template is entirely deterministic and accounts for most of the structure of the primes. Only the open positions — the ones not eliminated by small primes — contain any uncertainty.

The architecture of the skyline is therefore a deterministic sieve made visible. Width encodes divisibility. Height encodes factorization depth. Activation marks the entry of new layers. Coverage eliminates predictable fractions of the integers. Escape identifies the primes. Together, these elements form a geometric system in which the classical results of number theory arise naturally from the structure of the integers themselves.

3. Deterministic Structure of the Primes

Prime numbers have always been treated as if they were sprinkled unpredictably along the number line. The skyline overturns this intuition. In the skyline, primality is not a mysterious property but a geometric condition: a number is prime if and only if it escapes every coverage layer created by smaller primes. This escape principle is the foundation of the deterministic structure of the primes.

Every prime p creates a coverage layer that eliminates exactly one out of every p integers. These layers act independently, and their combined effect is multiplicative. After accounting for all primes up to p, the fraction of integers that remain uncovered is the escape density, the product of the escape fractions of each layer. This density is not a heuristic or an approximation; it is the exact geometric measure of how much space remains open after all smaller primes have claimed their territory.

A number escapes all coverage layers if and only if its least prime factor is itself. In other words, escape is primality. This identification is the key insight: primes are not special because they possess some intrinsic randomness, but because they are the survivors of a deterministic geometric process. The skyline makes this process visible. Primes appear as thin spires rising above the wide blocks of composite numbers, marking the positions that avoided every attempt at coverage.

This escape mechanism immediately explains why primes become rarer as numbers grow. Each new prime introduces a new coverage layer that eliminates a fixed fraction of the remaining space. The cumulative effect of these layers is the geometric form of the Prime Number Theorem: the number of primes up to N is approximately N times the escape density at $\sqrt{N}$. This density shrinks slowly but inexorably, producing the familiar thinning of primes. The skyline turns the Prime Number Theorem from an analytic statement into a geometric inevitability.

The same mechanism explains the distribution of smooth numbers. A number is y-smooth if all its prime factors are at most y. In the skyline, smooth numbers are the columns that can be peeled repeatedly by small primes. The peeling process mirrors the Dickman function, which describes the density of smooth numbers. What appears analytically as a differential-delay equation becomes geometrically the repeated removal of width layers. Smoothness and roughness are dual aspects of the same coverage product.

Other classical results fall out just as naturally. The density of squarefree numbers, $6/\pi^2$, is simply the probability that a number escapes the second layer of coverage by each prime. The Chebyshev function $\theta(x)$, the sum of the logarithms of primes up to x, becomes the cumulative height of the escape spires. Its linear growth reflects the fact that escape height accumulates at a constant rate.

The skyline therefore reveals a deterministic structure underlying the primes. Their distribution is governed by the geometry of coverage and escape, not by randomness. The irregularity of primes is not a sign of chaos but the natural outcome of independent layers acting on a rigid template. The skyline makes this structure explicit, showing that the primes are exactly where they must be.

4. Patterns in the Primes

Prime patterns—twin primes, triplets, quadruplets, and the larger constellations—have always been treated as deep mysteries. They appear irregularly, with no obvious structure, and their infinitude remains unproved in classical number theory. The Factor Skyline changes the perspective. In the skyline, prime patterns are not accidents or anomalies. They are the natural outcome of the geometry of coverage and escape.

The key idea is admissibility. A pattern of offsets, such as (0, 2) for twin primes or (0, 2, 6) for prime triplets, is admissible if no prime divides all of the positions simultaneously. This condition is simple, but in the skyline it becomes a geometric guarantee: if a pattern is admissible, then no coverage layer can eliminate all of its positions at once. Every prime p covers exactly one residue class modulo p, and admissibility ensures that the pattern never collapses entirely into that forbidden class. This is the geometric form of coverage protection.

Coverage protection is the reason prime constellations survive indefinitely. Each prime p removes a fraction 1/p of all possible placements of the pattern, but it never removes all of them. The surviving fraction is the survival factor, the proportion of placements that remain open after accounting for p. These survival factors multiply across primes, producing the Hardy–Littlewood constants. In the skyline, these constants are not mysterious analytic objects. They are the geometric weights that measure how much of the template remains open to a given pattern.

This geometric viewpoint also explains why some patterns are more common than others. Twin primes survive each coverage layer with a certain probability; triplets survive with a smaller one; quadruplets with a smaller one still. The skyline makes these differences visible. Each pattern has a characteristic footprint in the FS-x coordinate system, a fixed sequence of horizontal gaps between its escape points. These footprints are invariant: every sufficiently large instance of the pattern has the same geometric shape. This invariance is not a conjecture but a structural fact of the skyline.

The skyline therefore provides a unified explanation for prime patterns. Admissibility ensures that a pattern is never fully eliminated. Coverage protection guarantees that it survives every layer. The survival factors determine its frequency. The FS-x footprint encodes its geometric shape. Together, these elements show that prime

constellations are not rare coincidences but deterministic features of the escape geometry.

The classical conjectures about prime patterns—twin primes, prime triplets, prime k-tuples—are reinterpreted in this framework. Their infinitude is not a question of randomness but of whether the escape process continues to produce open positions at every scale. The skyline does not prove these conjectures, but it clarifies what a proof must address: the thin slice of uncertainty that remains after the template has done its work. This is the parity barrier in its geometric form.

Prime patterns, in the skyline, are not mysteries. They are the shadows cast by the deterministic structure of the integers, revealed by the geometry of coverage and escape.

5. Correlations and Randomness

The integers contain many kinds of statistical behavior. Some functions, like the Möbius function, behave almost like perfect randomness. Others, like the divisor function or $\omega(n)$, show persistent correlations. Prime gaps fluctuate in ways that look noisy but follow precise statistical laws. Even the zeros of the Riemann zeta function exhibit the same repulsion patterns seen in quantum systems.

Classically, these phenomena are treated separately. Each requires its own analytic machinery, its own conjectures, its own heuristics. The Factor Skyline reveals that they all come from a single source: the interaction between shared and independent coverage layers.

Every prime p creates a coverage layer that eliminates exactly one residue class modulo p. When comparing two numbers n and n+h, the layer p behaves in one of two ways. If p divides h, then n and n+h fall into the same residue class modulo p, and the layer acts on them together. If p does not divide h, then the layer acts independently on the two positions. This simple dichotomy—shared versus independent—determines the entire correlation structure of multiplicative functions.

For primes themselves, this leads directly to the Hardy–Littlewood pair correlation formula. When h has no small prime factors, almost all layers act independently, and the correlation is weak. When h is divisible by small primes, those layers act jointly, amplifying the chance that both positions escape together. This amplification produces the familiar constants: twin primes are more common than random chance would suggest, sexy primes even more so, and larger constellations follow the same pattern. The skyline makes this amplification visible: shared layers carve out synchronized escape channels.

For the Möbius function, the story is different. The Möbius value of a squarefree number depends on the parity of its prime factors. Parity flips are independent across layers unless p divides h, and the contribution from shared layers decays rapidly. As a result, the correlation between $\mu(n)$ and $\mu(n+h)$ tends to zero. This is the geometric form of the Chowla conjecture: the Möbius function decorrelates because almost all layers act independently.

The divisor function and $\omega(n)$ behave differently still. These functions count prime factors rather than detecting escape. Shared layers inflate both values simultaneously, creating persistent positive correlations. Even when h is large, any prime dividing h creates a shared layer that contributes to the correlation. The skyline explains why these correlations never fully disappear: the shared layers never go away.

At the highest level, the skyline even explains the behavior of the zeta zeros. The explicit formula links the zeros to oscillations generated by the primes. These oscillations are built from the logarithms of primes, which are

incommensurate. Their lack of alignment produces level repulsion, the hallmark of GUE statistics. In the skyline, this repulsion is the spectral analogue of independent coverage layers: the harmonics cannot synchronize, so the zeros cannot cluster.

All of these phenomena—prime pairs, Möbius cancellation, divisor correlations, ω(n) persistence, and GUE statistics—are manifestations of the same mechanism. Shared layers create structure; independent layers create randomness. The skyline unifies these behaviors into a single framework. What once required separate theories now emerges from the geometry of coverage and escape.

The result is a four-level hierarchy of correlations. Escape correlations govern primes. Parity correlations govern Möbius. Branching correlations govern divisor functions. Spectral correlations govern the zeros. Each level builds on the one below it, and all of them arise from the same architectural principle.

In the skyline, randomness is not a primitive concept. It is the shadow cast by independence.

6. The Randomness Hierarchy

One of the most striking features of the integers is the mixture of structure and randomness they seem to contain. Some patterns are rigid and predictable. Others look noisy and irregular. Prime gaps fluctuate. The Möbius function behaves like a coin flip. The zeros of the zeta function follow the same statistics as quantum energy levels. Classical number theory treats these phenomena separately, each with its own heuristics and conjectures.

The Factor Skyline reveals that all of these behaviors arise from a single architectural principle: the interaction between the deterministic template created by small primes and the thin layer of uncertainty that remains after the template has done its work. This interaction produces a hierarchy of randomness with three distinct levels.

The first level is the template, the periodic backbone created by the coverage layers of small primes. The template determines which positions are guaranteed composite and which remain open. It is completely rigid: once the primorial p# is fixed, the pattern of open and covered positions repeats forever. This rigidity is not approximate or statistical; it is exact. The template contains no randomness at all. It is the deterministic structure that underlies the entire skyline.

The second level is the escape layer, the set of open positions where the template does not decide the outcome. These positions are the only places where primes can occur. The escape process is deterministic, but because the coverage layers act independently, the pattern of escapes exhibits statistical behavior that resembles randomness. The skyline quantifies this behavior: only about 27 percent of the dx sequence is genuinely uncertain, and this uncertainty is concentrated entirely in the escape layer. The result is a pseudo-random process with sub-Poisson variance. Primes are less variable than a random sequence because the template imposes deterministic spacing, but they still fluctuate in ways that mimic filtered noise.

The third level is the spectral layer, the domain of the zeta zeros. The explicit formula links the zeros to oscillations generated by the primes. These oscillations are built from the logarithms of primes, which are incommensurate. Their lack of alignment produces level repulsion, the hallmark of GUE statistics. In the skyline, this repulsion is the spectral analogue of independent coverage layers: the harmonics cannot synchronize, so the zeros cannot cluster. The spectral layer therefore exhibits the highest possible entropy consistent with the constraints of the system.

These three layers—template, escape, and spectral—form the randomness hierarchy. Each layer has its own characteristic behavior. The template is rigid and periodic. The escape layer is pseudo-random and sub-Poisson. The spectral layer is maximally random and GUE-distributed. The hierarchy is not imposed from outside; it emerges naturally from the geometry of coverage and escape.

The skyline also explains the mixing behavior of different arithmetic functions. The divisor function and ω(n) are influenced by shared layers, so their correlations persist. The Möbius function depends on parity, which flips independently across layers, so its correlations vanish. The dx sequence is a mixture: 73 percent rigid, 27 percent mixing. These differences are not mysterious. They reflect how each function interacts with the randomness hierarchy.

In the skyline, randomness is not a primitive concept. It is the residue left after the deterministic structure has been accounted for. The hierarchy shows that the integers contain exactly as much randomness as the architecture allows—and no more.

7. Information Theory of the Skyline

One of the most surprising discoveries in the Factor Skyline is that primality, despite its apparent irregularity, contains very little true uncertainty. Most of what determines whether a number is prime is already fixed by the deterministic structure created by small primes. Only a thin residue of information remains undecided, and that residue is the exact source of the difficulty behind the major open problems in number theory.

Information theory gives us a way to measure uncertainty. Instead of asking whether primes are "random," we ask a sharper question: *How many yes/no questions would you need, on average, to determine whether a number is prime?* The skyline provides a precise answer.

The first component is the **template**, the periodic pattern of open and covered positions modulo the primorial p#. The template eliminates all numbers that are divisible by small primes. This elimination is not probabilistic; it is exact. Once the template is fixed, every number falls into a predetermined category: either it is guaranteed composite, or it remains open. The template therefore removes a large fraction of the uncertainty about primality. In fact, depending on how far the template is extended, it removes between 50 and 84 percent of the total information.

This means that most of the structure of the primes is already encoded in the periodic backbone of the integers. The apparent irregularity of primes does not come from the small primes; it comes from what happens after their work is done.

The second component is the **escape layer**, the set of open positions where the template does not decide the outcome. These positions are the only places where primes can occur. The skyline quantifies the uncertainty in this layer: the escape process carries about 0.26 bits of entropy per integer. This is the irreducible core of primality—the part that cannot be predicted from the template alone.

This number, 0.26 bits per integer, is not arbitrary. It is the information-theoretic form of the parity barrier. The barrier arises because the template can determine whether a number has at least one small prime factor, but it cannot distinguish between numbers with exactly one prime factor (the primes) and numbers with three or more prime factors. That distinction requires information that the coverage layers cannot supply. The skyline shows that this missing information is precisely the 0.26 bits per integer that remain after the template is accounted for.

The third component is **activation entropy**, the uncertainty associated with the entry of new coverage layers at

prime squares. Each prime p activates at p^2, and the probability that a random number falls into the new coverage class is 1/p. The entropy contributed by each layer is the binary entropy of this probability, and the total activation entropy is the sum of these contributions. This sum converges, reflecting the fact that activation events become increasingly rare.

The final component is **spectral entropy**, the entropy of the zeta zeros. The explicit formula links the zeros to oscillations generated by the primes. These oscillations are built from the logarithms of primes, which are incommensurate. Their lack of alignment produces the maximum possible entropy consistent with the constraints of the system. In the skyline, the spectral layer is the highest-entropy layer, the one that exhibits GUE statistics. Together, these components form the information-theoretic structure of the skyline. The template is rigid and carries no entropy. The escape layer carries a small but irreducible amount of entropy. The spectral layer carries the maximum entropy allowed by the architecture. This structure mirrors the randomness hierarchy and the ergodic behavior of the system.

The skyline therefore provides a precise answer to the question of how much information is contained in primality. Most of it is deterministic. A small but crucial part is genuinely uncertain. And that uncertainty is exactly the barrier that separates what the architecture can prove from what remains open.

8. Dynamics of the Skyline

The Factor Skyline is not just a static picture of the integers. It is a dynamical system. As the skyline grows, its structure evolves in predictable ways. Some components remain rigid forever. Others mix, fluctuate, and settle into statistical equilibrium. Still others exhibit the highest level of chaotic behavior permitted by the architecture. These dynamical properties are not imposed from outside; they emerge naturally from the geometry of coverage and escape.

The most rigid component is the template, the periodic backbone created by the coverage layers of small primes. Once the primorial p# is fixed, the pattern of open and covered positions repeats indefinitely. This periodicity makes the template trivially ergodic: every residue class appears with the correct frequency, but the system does not mix. The future is entirely determined by the past. The template is the deterministic skeleton of the skyline.

Above the template lies the escape layer, the domain of primality. The escape process is deterministic, but because the coverage layers act independently, the pattern of escapes exhibits statistical behavior. The skyline shows that the escape process is ergodic: the density of primes stabilizes, and the system explores its state space in a representative way. But it is not fully mixing. The mixing rate depends on the distribution of prime gaps, and under the Riemann Hypothesis it decays like the inverse square root of N. The escape layer is therefore a hybrid: deterministic in mechanism, statistical in behavior.

The divisor function and $\omega(n)$ reveal a different kind of dynamical behavior. These functions count prime factors rather than detecting escape. Shared coverage layers inflate both values simultaneously, creating persistent correlations. As a result, $\omega(n)$ is ergodic but not mixing. Its long-range correlations never fully disappear. The skyline explains this persistence: the shared layers never vanish, and their influence accumulates over time.

The Möbius function behaves differently still. Its value depends on the parity of the number of prime factors, and parity flips independently across layers unless p divides the offset. This independence destroys persistent correlations. The Möbius function is conjecturally mixing: its correlations decay to zero. In the skyline, this mixing is the dynamical counterpart of the parity cancellation that drives the Chowla conjecture.

The dx sequence, which encodes the widths of the columns, exhibits a mixture of behaviors. About 73 percent of its values are determined by the template, and the remaining 27 percent come from the escape layer. This mixture produces a quasi-ergodic system: asymptotically rigid but permanently influenced by the thin layer of uncertainty. The skyline quantifies this mixture precisely, showing how the deterministic and stochastic components interact.

At the highest level lies the spectral layer, the domain of the zeta zeros. The explicit formula links the zeros to oscillations generated by the primes. These oscillations are built from the logarithms of primes, which are incommensurate. Their lack of alignment produces level repulsion, the hallmark of GUE statistics. In dynamical terms, the spectral layer is maximally mixing: it exhibits the highest entropy and the strongest form of chaotic behavior permitted by the architecture.

These layers—template, escape, branching, and spectral—form a dynamical hierarchy. Each layer has its own characteristic behavior. The template is rigid. The escape layer is ergodic. The branching layer is non-mixing. The spectral layer is GUE-mixing. The hierarchy mirrors the randomness hierarchy and the information-theoretic structure of the skyline.

The skyline therefore reveals that the integers are not just a static sequence. They are a dynamical system with distinct regimes of behavior, each governed by the geometry of coverage and escape. The interplay between rigidity and randomness, between periodicity and mixing, is not accidental. It is the natural consequence of the architecture of the integers.

9. Universality

One of the most striking outcomes of the Factor Skyline is that its architecture is not unique to the integers. The skyline is built from four simple axioms: a countable set of generators with unique factorization, activation at the square of each generator, independence of coverage layers, and escape defined as the complement of coverage. These axioms are minimal, but they are powerful. Any number system that satisfies them exhibits the same structural behavior as the integers.

This means that the skyline is not just a picture of $\mathbb{Z}$. It is a universal architecture for multiplicative systems. The integers are simply one instance of a broader class.

The first universality class is the logarithmic class, where the escape density shrinks like a constant divided by the logarithm of x. The integers belong to this class. In these systems, primes thin out slowly, and the prime-counting function grows like x divided by log x. The skyline shows that this behavior is not a special property of $\mathbb{Z}$ but a structural consequence of the axioms. Any system with generators that grow like the primes will exhibit the same thinning.

The second universality class is the constant class, where the escape density approaches a positive constant. Polynomial rings over finite fields, such as $F_q[x]$, belong to this class. In these systems, the analogue of primes occurs with constant density, and the prime-counting function grows linearly. The skyline shows that this behavior is the natural outcome of having generators that grow linearly rather than multiplicatively. The architecture is the same; only the scale changes.

The third universality class is the degenerate class, where the escape density collapses rapidly to zero. In these systems, only finitely many escapes occur. The skyline shows that this collapse is the inevitable result of

generators that grow too quickly. The coverage layers overwhelm the system, leaving no room for escape.

Across all three classes, the skyline identifies twelve structural features that are universal. These include the product form of the escape density, the periodicity of the template, the geometric form of the prime-counting function, the normal distribution of ω(n), the cancellation of the Möbius function, the density of squarefree elements, the persistence of admissible patterns, the protection of constellations, the conservation law for escape height, the structure of the explicit formula, the sub-Poisson variance of escapes, and the existence of epochs. These features do not depend on the specific values of the generators. They follow directly from the axioms.

Even the GUE statistics of the zeta zeros have a universal explanation. The explicit formula links the zeros to oscillations generated by the logarithms of the generators. When these logarithms are incommensurate, the oscillations cannot align, and level repulsion emerges. This repulsion is the hallmark of GUE statistics. The skyline shows that this behavior is not unique to the Riemann zeta function. It is the natural outcome of independence and incommensurability.

The universality criterion is simple: a phenomenon is universal if it can be proved from the four axioms. If it requires the specific values of the generators, it is system-specific. This criterion provides a clean separation between what is structural and what is accidental. The skyline therefore offers a new way to understand the landscape of number systems: not as isolated objects but as instances of a single architectural framework.

The parity barrier is universal as well. The skyline shows that every system in every universality class contains a thin layer of irreducible uncertainty. This uncertainty is the information-theoretic gap between knowing that a number is open and knowing whether it escapes. The barrier is not a quirk of the integers. It is a structural limit that arises in every system governed by the axioms.

Universality reveals the true power of the skyline. It is not merely a visualization of the integers. It is a blueprint for the architecture of multiplicative systems. The integers are one example. The skyline is the theory.

10. Meta-Structure

The Factor Skyline is built from simple ingredients: the least prime factor, the geometry of width and height, the activation of coverage layers, and the escape mechanism that identifies primes. But beneath these ingredients lies a deeper structure. The skyline is not just a model of the integers; it is an architecture with its own internal logic, its own hierarchy of concepts, and its own system of dualities. Understanding this meta-structure reveals why the skyline works, why it is minimal, and why it is inevitable.

The foundation of the skyline is the least prime factor. From this single primitive, all other concepts emerge. Width is the least prime factor itself. Height is the quotient after removing it. Coverage is the action of width layers on the integers. Activation is the moment when a new width layer enters the system. Escape is the complement of coverage. This chain of derivations shows that the skyline does not require a large set of assumptions. It requires only the existence of a total order, multiplication, and unique factorization. These are the minimal axioms needed to define the least prime factor, and therefore the minimal axioms needed to build the skyline.

From this primitive foundation, the skyline generates a hierarchy of structures. At the first level are the basic geometric objects: widths, heights, escape density, and templates. These are the building blocks of the architecture. At the second level are the classical results of number theory: the Prime Number Theorem, the

Chebyshev functions, the persistence of admissible patterns, and the protection of constellations. These results emerge naturally from the first-order structures. At the third level are the spectral phenomena: the explicit formula, the distribution of zeros, and the appearance of GUE statistics. These phenomena arise from the interaction between the escape layer and the harmonic structure of the primes.

This hierarchy is not arbitrary. It reflects the way the skyline organizes information. Each level builds on the one below it, and each level reveals new behavior that cannot be seen at lower levels. The skyline therefore provides a natural stratification of number theory, one that mirrors the stratification of the integers themselves.

The skyline also contains a system of dualities. The first is the duality between coverage and escape. Coverage eliminates composite numbers; escape identifies primes. These two processes partition the integers. The second is the duality between width layers and the spectral layer. Width layers create the geometric structure of the skyline; the spectral layer encodes the oscillatory behavior of the zeta zeros. The explicit formula bridges these two domains. The third is the duality between rigidity and randomness. The template is rigid; the spectral layer is maximally random. Fourier analysis connects these extremes.

These dualities reveal that the skyline is self-referential. The structure of the architecture mirrors the structure of the integers. The hierarchy of concepts mirrors the hierarchy of results. The limits of the architecture mirror the limits of what can be proved. Even the parity barrier has a meta-structural analogue: just as coverage cannot determine escape, the architecture cannot determine certain global properties of the integers. The skyline therefore contains within itself the explanation for its own limitations.

The meta-structure of the skyline shows that the framework is not a collection of independent ideas. It is a coherent system in which every concept, every theorem, and every phenomenon arises from a small set of primitives. The skyline is not just a way of visualizing the integers. It is a way of understanding the architecture of number theory itself.

11. Meta-Mathematics

The Factor Skyline is built from simple primitives, but beneath its geometric and statistical structure lies a deeper layer: the meta-mathematics of the architecture itself. This layer concerns what the skyline can compute, what it can compress, what it can prove, and where its limits lie. These questions are not external to the framework. They arise naturally from the way the skyline organizes information and from the constraints imposed by its primitives.

The first meta-mathematical feature of the skyline is its computability. The skyline up to N can be constructed in essentially the same time it takes to compute the least prime factor of each number up to N. Using classical sieve methods, this requires O(N log log N) operations. The template statistics, which depend only on the primorial p#, can be computed in time proportional to the number of primes involved. The escape layer, which identifies primes, is as hard to compute as primality itself. This alignment between the computational complexity of the skyline and the complexity of the integers is not accidental. It reflects the fact that the skyline is a faithful geometric representation of the multiplicative structure of $\mathbb{Z}$.

The second feature is pseudo-randomness. The skyline reveals that the escape layer contains about 0.26 bits of entropy per integer. This entropy is genuine: it cannot be predicted from the template. But the Kolmogorov complexity of the skyline up to N is only O(log N). This means that the skyline is highly compressible. The deterministic structure dominates, and the pseudo-random component is thin. The skyline therefore exhibits a striking contrast: it contains a large amount of statistical randomness but a small amount of algorithmic

randomness. This contrast is the meta-mathematical form of the parity barrier.

The third feature concerns logical limits. The skyline shows that the parity barrier is not a form of Gödel incompleteness. It does not arise from the limitations of formal logic. It arises from the limitations of coverage-based arguments. Coverage can determine whether a number has a small prime factor, but it cannot determine whether a number has exactly one prime factor or three or more. This distinction requires information that the coverage layers cannot supply. The skyline therefore identifies the methodological boundary of sieve-based proofs. It explains why certain conjectures remain open: not because they are unprovable, but because the available methods cannot cross the information barrier.

The skyline also defines a first-order theory, T_FS, built from the four universality axioms. This theory has models beyond the integers, including Gaussian integers and polynomial rings over finite fields. The theory is incomplete: different models satisfy the same axioms but differ in their specific behaviors. This incompleteness is not a weakness. It reflects the universality of the architecture. The skyline captures the structural features common to all multiplicative systems with unique factorization, but it does not determine the specific values of the generators. Those values are system-specific.

The final meta-mathematical feature is the computational–mathematical mirror. The skyline reveals a parallel between the computational complexity of different layers and the difficulty of proving theorems about them. The template is trivial to compute and trivial to prove. Density results require logarithmic computation and elementary proofs. Escape requires near-linear computation and deep theorems. The spectral layer requires unknown computation and lies at the frontier of mathematics. This mirror shows that the difficulty of computation and the difficulty of proof are aligned. The skyline therefore provides a unified perspective on the landscape of number theory: the computational and logical challenges arise from the same structural features.

The meta-mathematics of the skyline shows that the architecture is not just a geometric or statistical framework. It is a computational and logical system with precise limits and strengths. It explains what the skyline can compute, what it can compress, what it can prove, and where it must stop. These limits are not arbitrary. They are the natural consequences of the architecture of the integers.

12. The Unified Meta-Architecture

The Factor Skyline begins as a geometric reinterpretation of the integers, but by the time the full framework is assembled, something remarkable becomes clear: the architecture of the skyline mirrors the architecture of number theory itself. Every structural feature of the skyline has a meta-structural analogue. Every mathematical phenomenon has a computational counterpart. Every limit in the skyline corresponds to a limit in what can be proved. The framework is self-similar at every scale.

This self-similarity is the unified meta-architecture.

At the foundation of the skyline lies the least prime factor. From this single primitive, the entire geometric structure emerges: widths, heights, coverage layers, activation points, escape density, templates, and epochs. At the meta-level, the same primitive generates the computational and logical structure: the sieve program, the universality axioms, the emergence hierarchy, and the limits of coverage-based reasoning. The skyline therefore contains a dual architecture: one at the level of the integers, and one at the level of the theory itself.

The first component of the meta-architecture is the meta-correspondence. Every concept in the skyline has a

meta-mathematical counterpart. Coverage corresponds to compression: both eliminate redundancy. Escape corresponds to incompressibility: both identify the irreducible core. Activation corresponds to the introduction of new information. Templates correspond to periodic structure in logic and computation. The explicit formula corresponds to the bridge between geometry and spectrum. These correspondences are not analogies. They are structural identities. The skyline is a compression algorithm for number theory, and number theory is the expansion of that algorithm.

The second component is the boundary between architecture and logic. The skyline reveals that the parity barrier — the thin layer of irreducible uncertainty in the escape process — has four equivalent descriptions. It is an information-theoretic boundary: about 0.26 bits per integer that the template cannot supply. It is a computational boundary: the escape layer is as hard to compute as primality. It is a logical boundary: the difference between knowing that a pattern persists (a Σ_1 fact) and knowing that it persists infinitely often (a Π_2 fact). And it is a universal boundary: it appears in every system governed by the universality axioms. These four descriptions are the same phenomenon seen from different angles. The skyline therefore identifies the exact point where the architecture stops and where new ideas are required.

The third component is the emergence hierarchy, which appears both in the skyline and in its meta-structure. At the object level, the hierarchy runs from widths and heights to the Prime Number Theorem, and from there to the explicit formula and the distribution of zeros. At the meta-level, the hierarchy runs from the least prime factor to the universality axioms, and from there to the computational-mathematical mirror and the limits of provability. Each level builds on the one below it. Each level reveals new behavior that cannot be seen at lower levels. The skyline therefore provides a unified view of emergence in mathematics: simple primitives give rise to complex structures through deterministic rules.

The fourth component is the duality structure. Coverage and escape form one duality. Width layers and spectral layers form another. Rigidity and randomness form a third. These dualities appear at the object level, but they also appear at the meta-level: compression and incompressibility, decidability and undecidability, structure and noise. The skyline therefore contains a nested system of dualities, each reflecting the others.

The unified meta-architecture shows that the skyline is not just a framework for understanding the integers. It is a framework for understanding the structure of mathematical theories. It explains why certain results are easy, why others are hard, and why some remain open. It identifies the minimal assumptions needed to build the architecture and the maximal consequences that follow from those assumptions. It reveals the deep symmetry between the integers and the theory that describes them.

In the end, the skyline is not just a new way of seeing the integers. It is a new way of seeing mathematical structure itself. The architecture and the meta-architecture are reflections of the same underlying system. The skyline is the integers, and the integers are the skyline.

13. Appendix

The appendix collects the numerical data, computational procedures, and reproducibility materials that support the results presented in the main text. The Factor Skyline is a geometric and conceptual framework, but it is also an empirical one: many of its structural claims can be illustrated, verified, or explored through explicit computation. The appendix provides the raw material for that exploration.

The first component is the set of numerical tables (on GitHub). These tables record the skyline coordinates of

small integers, the convergence of escape density, the statistics of primorial templates, and the verification of classical results such as the Prime Number Theorem, the conservation law for escape height, and the Dickman function. They also include data on squarefree numbers, Möbius values, divisor averages, the distribution of ω(n), and the FS-x footprints of prime constellations. These tables are not proofs, but they provide concrete illustrations of the architecture. They show how the skyline behaves at small scales and how its predictions align with classical number theory.

The second component is the set of correlation and randomness tables (on GitHub). These include prime pair counts for various offsets, Möbius correlations, divisor correlations, ω(n) correlations, and the normalized spacings of the first several zeta zeros. They also include measurements of sub-Poisson variance in prime counts and autocorrelation data for the dx sequence. These tables demonstrate the statistical behavior of the skyline and provide empirical support for the correlation hierarchy and the randomness hierarchy.

The third component is the reproducibility layer (on GitHub). All computations in the appendix were performed using the scripts in the project repository. The skyline coordinates, template statistics, escape density calculations, and correlation measurements can be reproduced exactly using the provided code. This reproducibility is essential. The skyline is not a heuristic model; it is a deterministic architecture. The appendix ensures that every numerical claim in the monograph can be verified independently.

The fourth component is the reference list, which situates the skyline within the broader context of number theory. The references include classical works on multiplicative number theory, analytic number theory, prime distribution, random matrix theory, and information theory. They also include the original Factor Skyline preprint. These references provide the theoretical background for the results in the monograph and connect the skyline to the existing literature.

The appendix therefore serves three purposes. It provides concrete examples of the skyline's geometry. It documents the statistical behavior of the escape process and its associated functions. And it ensures that the entire framework is reproducible, transparent, and grounded in explicit computation. The skyline is a conceptual architecture, but it is also a computational one. The appendix is the bridge between the two.

14. Glossary

This glossary collects the core terms of the Factor Skyline. Each entry is written to clarify not only the definition but the role the concept plays in the architecture. The skyline is built from a small set of primitives, and nearly every phenomenon in the monograph can be traced back to these definitions.

Activation
The moment when a new width-p coverage layer enters the skyline, occurring at p^2. Below p^2, every multiple of p has a smaller prime factor; at p^2, p becomes the least prime factor for the first time. Activation marks the beginning of a new epoch.

Admissible pattern
A set of offsets (such as (0,2) for twin primes) that never collapses entirely into a single residue class modulo any prime. Admissibility guarantees that no coverage layer can eliminate all positions of the pattern, ensuring its survival.

Column

The geometric representation of an integer in the skyline. A column has width equal to its least prime factor and height equal to the quotient after dividing by that factor. Columns are placed consecutively to form the skyline.

Coverage
The action of a width-p layer eliminating all integers divisible by p. Each layer removes exactly 1/p of all positions. Coverage is deterministic and multiplicative across primes.

Coverage protection
The principle that no width-p layer can eliminate all positions of an admissible pattern. This protection ensures the persistence of prime constellations.

Epoch
The interval between consecutive prime squares, $[p_k^2, p_{k+1}^2)$. Within an epoch, no new width layers activate, so the coverage pattern is frozen. Epochs are the natural "time slices" of the skyline.

Escape
The event that a number avoids all coverage layers. A number escapes if and only if its least prime factor is itself. Escape is primality.

Escape density
The fraction of integers that escape all coverage layers up to p. It is the product of the escape fractions of each layer: $D(p) = \prod_{q \le p} (1 - 1/q)$. This density governs the distribution of primes.

FS-x coordinate
The horizontal position of a column in the skyline, defined as the cumulative sum of widths. Primes contribute width 1; composites contribute width equal to their least prime factor.

FS-x footprint
The invariant pattern of horizontal gaps between the escape points of a prime constellation. Each admissible pattern has a fixed footprint that appears in every sufficiently large instance.

Height
The vertical dimension of a column, equal to n divided by its least prime factor. Height encodes the depth of the factorization process.

Least prime factor (lpf)
The smallest prime dividing n. This single function generates the entire skyline architecture: widths, heights, coverage, activation, and escape.

Parity barrier
The thin layer of irreducible uncertainty in the escape process, about 0.26 bits per integer. It reflects the fact that coverage can detect small prime factors but cannot distinguish primes from numbers with three or more prime factors.

Primorial (p#)
The product of all primes up to p. The primorial defines the period of the template and the backbone of the skyline's periodic structure.

Quasi-ergodic
A system that is mostly rigid but contains a persistent mixing component. The dx sequence is quasi-ergodic: 73 percent determined by the template, 27 percent mixing from the escape layer.

Renormalization
The extension of the template from $p_k\#$ to $p_{k+1}\#$. Each extension introduces a new coverage layer and reduces the escape density. Renormalization describes the flow of the skyline as it grows.

Spectral layer
The domain of the zeta zeros. The spectral layer exhibits GUE statistics due to the incommensurability of prime logarithms. It is the highest-entropy layer of the skyline.

Sub-Poisson
A statistical regime where the variance of escape counts is smaller than the mean. Primes exhibit sub-Poisson variance because the template imposes deterministic spacing.

Survival factor
The fraction of placements of an admissible pattern that remain open after accounting for a prime p. The product of survival factors across primes yields the Hardy–Littlewood constants.

Template
The periodic pattern of open and covered positions modulo the primorial p#. The template is rigid and accounts for most of the structure of the primes.

Template persistence
The fact that the number of open placements of an admissible pattern grows without bound as the template extends. This persistence reflects the geometric inevitability of prime constellations.

Width
The horizontal dimension of a column, equal to the least prime factor of the number. Width encodes divisibility and determines the skyline's geometry.